

19CSE456 Neural Network and Deep Learning Laboratory

List of Experiments

Week #	Experiment Title
1	Introduction to the lab and Implementation of a simple Perceptron (Hardcoding)
2	Implementation of Perceptron for Logic Gates (Hardcoding, Sklearn, TF)
3	Implementation of Multilayer Perceptron for XOR Gate and other classification problems with ML toy datasets (Hardcoding & TF)
4	Implementation of MLP for Image Classification with MNIST dataset (Hardcoding & TF)
5	Activation Functions, Loss Functions, Optimizers (Hardcoding & TF)
6	Lab Evaluation 1 (based on topics covered from w1 to w5)
7	Convolution Neural Networks for Toy Datasets (MNIST & CIFAR)
8	Convolution Neural Networks for Image Classification (Oxford Pets, Tiny ImageNet, etc.)
9	Recurrent Neural Networks for Sentiment Analysis with IMDB Movie Reviews
10	Long Short Term Memory for Stock Prices (Yahoo Finance API)



List of Experiments

contd.

Week #	Experiment Title
11	Implementation of Autoencoders and Denoising Autoencoders (MNIST/CIFAR)
12	Boltzmann Machines (MNIST/CIFAR)
13	Restricted Boltzmann Machines (MNIST/CIFAR)
14	Hopfield Neural Networks (MNIST/CIFAR)
15	Lab Evaluation 2 (based on CNN, RNN, LSTM, and AEs)
16	Case Study Review (Phase 1)
17	Case Study Review (Phase 1)

Optimization Techniques in TensorFlow

TensorFlow offers various gradient-based optimizers to minimize loss functions efficiently:

Gradient-Based Optimizers

1. Stochastic Gradient Descent
`tf.keras.optimizers.SGD()`
2. Adam
`tf.keras.optimizers.Adam()`
3. RMSprop
`tf.keras.optimizers.RMSprop()`
4. Adagrad
`tf.keras.optimizers.Adagrad()`
5. Adadelta
`tf.keras.optimizers.Adadelta()`

Learning Rate Schedulers

1. Exponential Decay
`tf.keras.optimizers.schedules.ExponentialDecay()`
2. Piecewise Constant Decay
`tf.keras.optimizers.schedules.PiecewiseConstantDecay()`
3. Cosine Decay `tf.keras.optimizers.schedules.CosineDecay()`
4. Inverse Time Decay
`tf.keras.optimizers.schedules.InverseTimeDecay()`
5. Polynomial Decay
`tf.keras.optimizers.schedules.PolynomialDecay()`

Gradient-Based Optimizers

```
optimizers = {  
    "Batch Gradient Descent (BGD)": tf.keras.optimizers.SGD(learning_rate=0.01), # Full batch size  
    "Mini-Batch Gradient Descent (MBGD)": tf.keras.optimizers.SGD(learning_rate=0.01), # Mini-batch size (e.g., 32)  
    "SGD": tf.keras.optimizers.SGD(learning_rate=0.01), # Stochastic GD (batch_size = 1)  
    "SGD with Momentum": tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9),  
    "NAG (Nesterov)": tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=True),  
    "Adam": tf.keras.optimizers.Adam(learning_rate=0.001),  
    "RMSprop": tf.keras.optimizers.RMSprop(learning_rate=0.001),  
    "Adagrad": tf.keras.optimizers.Adagrad(learning_rate=0.01),  
    "Adadelat": tf.keras.optimizers.Adadelat(learning_rate=1.0),  
    "Nadam": tf.keras.optimizers.Nadam(learning_rate=0.001),}
```

Gradient-Based Optimizers

```
# Model Creation
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Dense(64, activation='relu', input_shape=(input_shape,)),  
    tf.keras.layers.Dense(32, activation='relu'),  
    tf.keras.layers.Dense(2, activation='softmax')  
])
```

```
# Compile model
```

```
model.compile(  
    optimizer=optimizer,  
    loss='categorical_crossentropy',  
    metrics=['accuracy']  
)
```

```
# Train model
```

```
model.fit(  
    X_train, y_train,  
    epochs=10,  
    batch_size=batch_size,  
    verbose=0,  
    callbacks=[custom_callback]  
)
```

Learning Rate Schedulers

Step Decay

```
lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate=0.01, decay_steps=10000, decay_rate=0.9, staircase=True )  
optimizer = tf.keras.optimizers.SGD(learning_rate=lr_schedule)
```

The learning rate is reduced at specific intervals or "steps" during training

Exponential Decay

```
lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate=0.01, decay_steps=1000, decay_rate=0.96, staircase=False )  
optimizer = tf.keras.optimizers.Adam(learning_rate=lr_schedule)
```

To reduce the learning rate of a model during training gradually

Piecewise Decay

```
lr_schedule = tf.keras.optimizers.schedules.PiecewiseConstantDecay(  
    boundaries=[5000, 10000], values=[0.01, 0.005, 0.001] )  
optimizer = tf.keras.optimizers.SGD(learning_rate=lr_schedule)
```

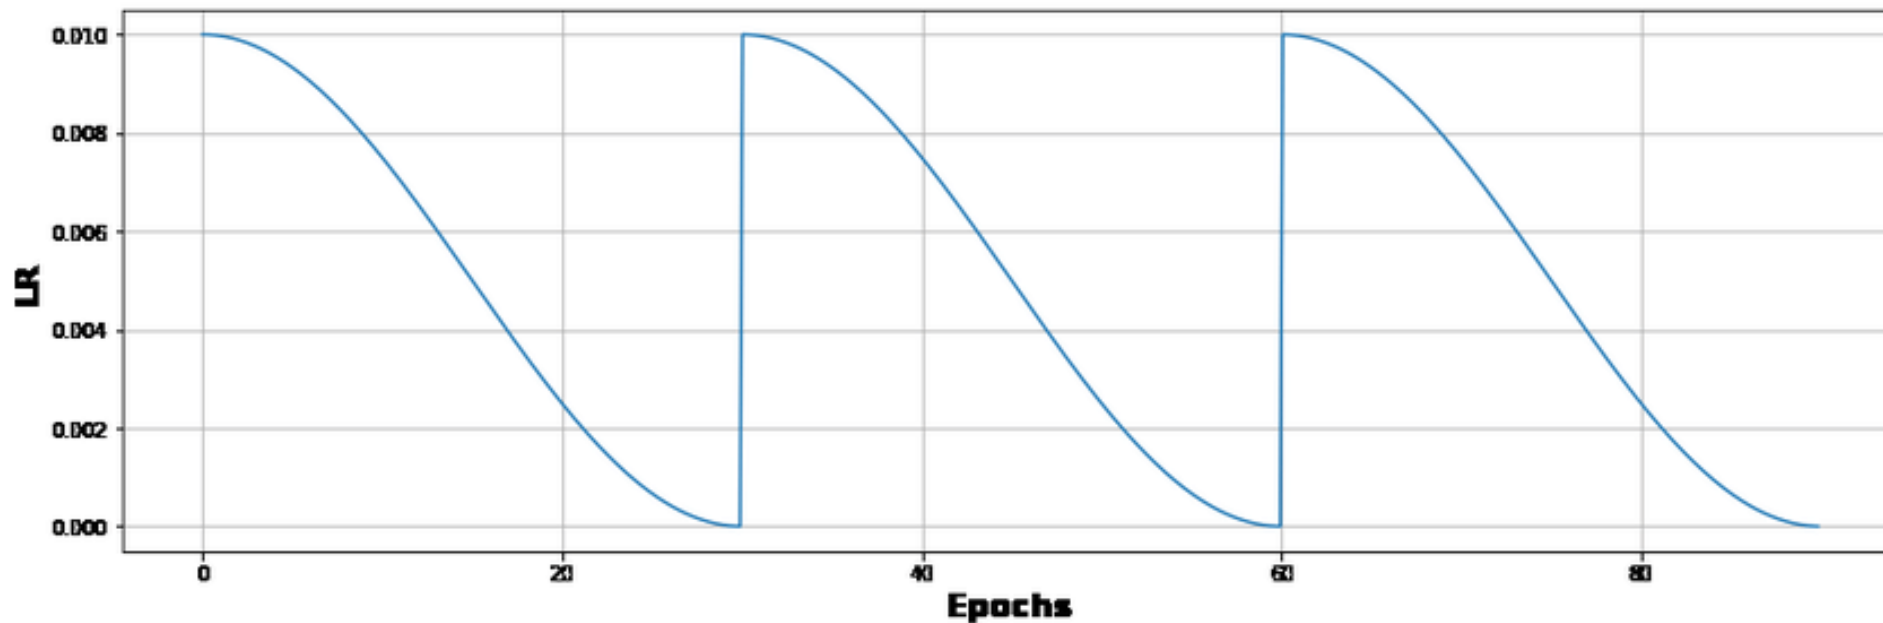
The learning rate remains constant within specified intervals (or "pieces") of training steps, but changes to a different value at the boundaries between these intervals

Learning Rate Schedulers

Cosine Decay

```
lr_schedule = tf.keras.optimizers.schedules.CosineDecay(  
    initial_learning_rate=0.01, decay_steps=10000, alpha=0.1 )  
optimizer = tf.keras.optimizers.Adam(learning_rate=lr_schedule)
```

Reduces the learning rate following a cosine function



Gradient-Based Optimizers

```
# Model Creation
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Dense(64, activation='relu', input_shape=(input_shape,)),  
    tf.keras.layers.Dense(32, activation='relu'),  
    tf.keras.layers.Dense(2, activation='softmax')  
])
```

```
# Compile model
```

```
model.compile(  
    optimizer=optimizer,  
    loss='categorical_crossentropy',  
    metrics=['accuracy']  
)
```

```
# Train model
```

```
model.fit(  
    X_train, y_train,  
    epochs=10,  
    batch_size=batch_size,  
    verbose=0,  
    callbacks=[custom_callback]  
)
```

Week 6 Exercises

1. Optimizers & Learning Rate Schedule in MLP Classifier

Objective: Implement and compare various optimizers and learning rate scheduling schemes in TensorFlow using the Titanic dataset.

Tasks:

- Load and preprocess the Titanic dataset.
- Implement a MLP classifier in TensorFlow.
 - Train the model using the following optimizers:
 - SGD (Stochastic)
 - Momentum
 - Adam
 - RMSprop
- Apply the following learning rate scheduling schemes:
 - Step Decay
 - Exponential Decay
 - Piecewise Constant Decay
 - Cosine Decay
- Compare the performance of the different optimizers and learning rate schedules by evaluating metrics such as accuracy and loss.